
PySpark Data Audit



PySparkAudit: PySpark Data Audit Library API

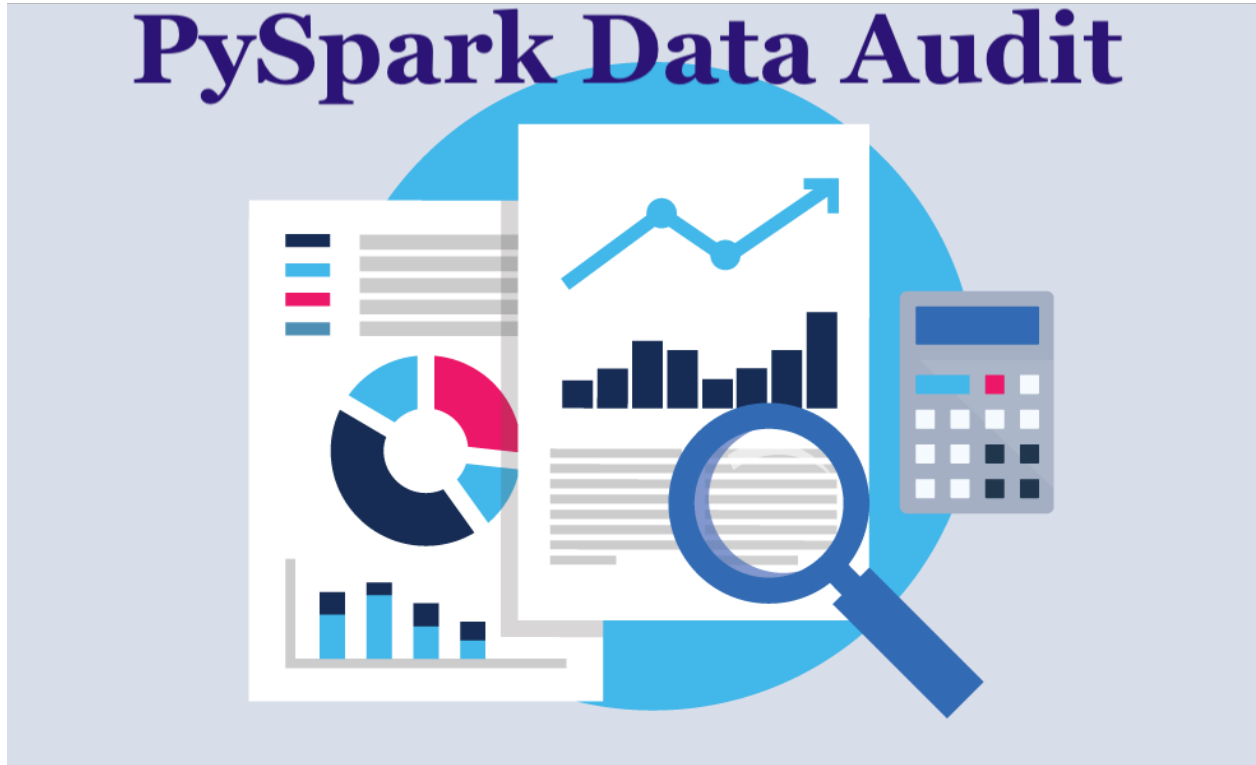
Wenqiang Feng and Yiming Xu

June 26, 2019

CONTENTS

1	Preface	3
1.1	About	3
1.1.1	About this API	3
1.1.2	About the author	3
1.2	Acknowledgement	4
1.3	Feedback and suggestions	4
2	How to Install	5
2.1	Install with <code>pip</code>	5
2.2	Install from Repo	5
2.2.1	Clone the Repository	5
2.2.2	Install	5
2.2.3	Uninstall	5
2.2.4	Test	6
3	Python Data Audit Functions	9
3.1	Basic Functions	9
3.1.1	<code>dtypes_class</code>	9
3.1.2	<code>missing_rate</code>	9
3.1.3	<code>zero_rate</code>	10
3.1.4	<code>feature_variance</code>	10
3.1.5	<code>freq_items_df</code>	11
3.1.6	<code>feature_len</code>	12
3.1.7	<code>correlation matrix</code>	12
3.2	Summary Functions	13
3.2.1	<code>numeric_summary</code>	13
3.2.2	<code>category_summary</code>	13
3.3	Auditing Function	14
3.3.1	<code>auditing</code>	14
4	Auditing Demos	15
4.1	Auditing in one function	15
4.2	Auditing function by function	17

5 Main Reference	23
Bibliography	25
Python Module Index	27
Index	29



Welcome to our **PySparkAudit: PySpark Data Audit Library API**! The PDF version can be downloaded from [HERE](#).

You can install the PySparkAudit from [PyPI](<https://pypi.org/project/PySparkAudit>):

```
pip install PySparkAudit
```


PREFACE

Chinese proverb

Good tools are prerequisite to the successful execution of a job. – old Chinese proverb

1.1 About

1.1.1 About this API

This document is the API book for our PyAudit: Python Data Audit Library [PyAudit] API. The PDF version can be downloaded from [HERE](#). **You may download and distribute it. Please be aware, however, that the note contains typos as well as inaccurate or incorrect description.**

The API assumes that the reader has a preliminary knowledge of `python` programing and `Linux`. And this document is generated automatically by using `sphinx`.

1.1.2 About the author

- **Wenqiang Feng**
 - Sr. Data Scientist and PhD in Mathematics
 - University of Tennessee at Knoxville
 - Webpage: <http://web.utk.edu/~wfeng1/>
 - Email: von198@gmail.com
- **Ming Chen**
 - Data Scientist and PhD in Genome Science and Technology
 - University of Tennessee at Knoxville

– Email: ming.chen0919@gmail.com

- **Biography**

Wenqiang Feng is Data Scientist within DST’s Applied Analytics Group. Dr. Feng’s responsibilities include providing DST clients with access to cutting-edge skills and technologies, including Big Data analytic solutions, advanced analytic and data enhancement techniques and modeling.

Dr. Feng has deep analytic expertise in data mining, analytic systems, machine learning algorithms, business intelligence, and applying Big Data tools to strategically solve industry problems in a cross-functional business. Before joining DST, Dr. Feng was an IMA Data Science Fellow at The Institute for Mathematics and its Applications (IMA) at the University of Minnesota. While there, he helped startup companies make marketing decisions based on deep predictive analytics.

Dr. Feng graduated from University of Tennessee, Knoxville, with Ph.D. in Computational Mathematics and Master’s degree in Statistics. He also holds Master’s degree in Computational Mathematics from Missouri University of Science and Technology (MST) and Master’s degree in Applied Mathematics from the University of Science and Technology of China (USTC).

- **Declaration**

The work of Wenqiang Feng was supported by the IMA, while working at IMA. However, any opinion, finding, and conclusions or recommendations expressed in this material are those of the author and do not necessarily reflect the views of the IMA, UTK and DST.

1.2 Acknowledgement

At here, Wenqiang Feng would like to thank **Weiyu Wang** at Missouri University of Science and Technology and **Jiangtao (Lotto) Xie** at Purdue University for the unit testing and valuable discussion.

1.3 Feedback and suggestions

Your comments and suggestions are highly appreciated. I am more than happy to receive corrections, suggestions or feedbacks through email (Wenqiang Feng: von198@gmail.com and Ming Chen: ming.chen0919@gmail.com) for improvements.

HOW TO INSTALL

2.1 Install with `pip`

You can install the `PyAudit` from [PyPI](<https://pypi.org/project/PyAudit>):

```
pip install PyAudit
```

2.2 Install from Repo

2.2.1 Clone the Repository

```
git clone https://github.com/runawayhorse001/PyAudit.git
```

2.2.2 Install

```
cd PyAudit
pip install -r requirements.txt
python setup.py install
```

2.2.3 Uninstall

```
pip uninstall statspy
```

2.2.4 Test

```
cd PyAudit/test
python test.py
```

test.py

Results:

```

feature  missing_rate
0        A           0.25
1        B           0.00
2        C           0.25
feature  zero_rate
0        A    0.333333
1        B    0.750000
2        C    0.000000
feature  feature_variance
0        A             1.0
1        B             0.5
2        C             1.0
  Age    Sex  ChestPain  RestBP  Chol  ...  Oldpeak  Slope  Ca
→  Thal  AHD
0   63   True    typical    145   233  ...    2.3     3  0.0
→  fixed  No
1   67   True  asymptomatic    160   286  ...    1.5     2  3.0
→  normal Yes
2   67   True  asymptomatic    120   229  ...    2.6     2  2.0
→reversible Yes
3   37   True   nonanginal    130   250  ...    3.5     3  0.0
→  normal  No
4   41  False   nontypical    130   204  ...    1.4     1  0.0
→  normal  No

[5 rows x 14 columns]
['Age', 'RestBP', 'Chol', 'Fbs', 'RestECG', 'MaxHR', 'ExAng', 'Oldpeak
→', 'Slope', 'Ca']
['ChestPain', 'Thal', 'AHD']
['Sex']
feature  dtypes
0        Age    int64
1        Sex     bool
2  ChestPain  object
3    RestBP    int64
4        Chol    int64
5        Fbs     int64
6    RestECG    int64
```

(continues on next page)

(continued from previous page)

```

7      MaxHR      int64
8      ExAng      int64
9      Oldpeak    float64
10     Slope      int64
11     Ca         float64
12     Thal       object
13     AHD        object
      feature      dtypes      class
0      Age        int64      numeric
1      Sex        bool       bool
2      ChestPain  object     category
3      RestBP     int64      numeric
4      Chol       int64      numeric
5      Fbs        int64      numeric
6      RestECG    int64      numeric
7      MaxHR      int64      numeric
8      ExAng      int64      numeric
9      Oldpeak    float64     numeric
10     Slope      int64      numeric
11     Ca         float64     numeric
12     Thal       object     category
13     AHD        object     category
      feature      missing_rate
0      Age        0.000000
1      Sex        0.000000
2      ChestPain  0.000000
3      RestBP     0.000000
4      Chol       0.000000
5      Fbs        0.000000
6      RestECG    0.000000
7      MaxHR      0.000000
8      ExAng      0.000000
9      Oldpeak    0.000000
10     Slope      0.000000
11     Ca         0.013201
12     Thal       0.006601
13     AHD        0.000000

Process finished with exit code 0

```


PYTHON DATA AUDIT FUNCTIONS

3.1 Basic Functions

3.1.1 dtypes_class

`PyAudit.basics.dtypes_class(df_in)`
numerical, categorical and bool name list in the DataFrame

Parameters `df_in` – input pandas DataFrame

Returns numerical, categorical and bool name list

Author Wenqiang Feng and Ming Chen

Email von198@gmail.com

```
>>> from PyAudit.basics import dtypes_class
>>> df = pd.read_csv('Heart.csv', dtype={'Sex': bool})
>>> (num_fields, cat_fields, bool_fields, data_types, type_class),
    => dtypes_class(df)
>>> num_fields
['Age', 'RestBP', 'Chol', 'Fbs', 'RestECG', 'MaxHR', 'ExAng',
 => 'Oldpeak', 'Slope', 'Ca']
```

3.1.2 missing_rate

`PyAudit.basics.missing_rate(df_in)`
calculate missing rate for each feature in the DataFrame

Parameters `df_in` – input pandas DataFrame

Returns missing rate

Author Wenqiang Feng and Ming Chen

Email von198@gmail.com

```
>>> import pandas as pd
>>> d = {'A': [1, 0, None, 3],
        'B': [1, 0, 0, 0],
        'C': ['a', None, 'c', 'd']}
>>> # create DataFrame
>>> df = pd.DataFrame(d)
>>> from PyAudit.basics import missing_rate
>>> missing_rate(df)
      feature  missing_rate
0          A           0.25
1          B           0.00
2          C           0.25
```

3.1.3 zero_rate

`PyAudit.basics.zero_rate(df_in)`

calculate the percentage of 0 value for each feature in the DataFrame

Parameters `df_in` – input pandas DataFrame

Returns zero rate

Author Wenqiang Feng and Ming Chen

Email von198@gmail.com

```
>>> import pandas as pd
>>> d = {'A': [1, 0, None, 3],
        'B': [1, 0, 0, 0],
        'C': ['a', None, 'c', 'd']}
>>> # create DataFrame
>>> df = pd.DataFrame(d)
>>> from PyAudit.basics import zero_rate
>>> zero_rate(df)
      feature  zero_rate
0          A   0.333333
1          B   0.750000
2          C   0.000000
```

3.1.4 feature_variance

`PyAudit.basics.feature_variance(df_in)`

calculate the variance for each feature

Parameters `df_in` – input pandas DataFrame

Returns feature variance

Author Wenqiang Feng and Ming Chen

Email von198@gmail.com

```
>>> import pandas as pd
>>> d = {'A': [1, 0, None, 3],
        'B': [1, 0, 0, 0],
        'C': ['a', None, 'c', 'd']}
>>> # create DataFrame
>>> df = pd.DataFrame(d)
>>> from PyAudit.basics import zero_rate
>>> zero_rate(df)
      feature  feature_variance
0          A                1.0
1          B                0.5
2          C                1.0
```

3.1.5 freq_items_df

`PyAudit.basics.freq_items_df(df_in, top_n=3)`

find out the top n values and the corresponding frequency for each feature

Parameters

- **df_in** – input pandas DataFrame
- **top_n** – the number of the top values

Returns top n values and the corresponding frequency for each feature

Author Wenqiang Feng and Ming Chen

Email von198@gmail.com

```
>>> d = {
>>>     'num': list('1223334444'),
>>>     'cat': list('wxyyyzzzz')
>>> }
>>> df = pd.DataFrame(d)
>>> df = df.astype({"num": int, "cat": object})
>>> print(freq_items_df(df, top_n=4))
      feature  top_items  top_freqs
0        num  [4, 3, 2, 1]  [4, 3, 2, 1]
1        cat  [z, y, x, w]  [4, 3, 2, 1]
```

3.1.6 feature_len

`PyAudit.basics.feature_len(df_in)`

find out the min and max length of values for each feature

Parameters `df_in` – input pandas DataFrame

Returns min and max length DataFrame

Author Wenqiang Feng and Ming Chen

Email von198@gmail.com

```
>>> d = {'A': [1, 0, None, 3],
>>>        'B': [1, 0, 0, 0],
>>>        'C': ['a', None, 'c', 'd']}
>>> # create DataFrame
>>> df = pd.DataFrame(d)
>>> print(df)
   A  B  C
0  1  1  a
1  0  0 None
2 NaN 0  c
3  3  0  d
>>> print(feature_len(df))
   feature  min_length  max_length
0         A           3           3
1         B           1           1
2         C           1           4
```

3.1.7 correlation matrix

`PyAudit.basics.corr_matrix(df_in, output_dir)`

generate correlation matrix for numerical dataframe

Parameters

- `df_in` – input pandas DataFrame
- `output_dir` – output path

Returns correlation matrix

Author Wenqiang Feng and Ming Chen

Email von198@gmail.com

```
>>> d = {'A': [1, 0, None, 3],
>>>        'B': [1, 0, 0, 0],
```

(continues on next page)

(continued from previous page)

```

>>>      'C': ['a', None, 'c', 'd']}
>>> # create DataFrame
>>> df = pd.DataFrame(d)
>>> print(corr_matrix(df))
           A          B
A  1.000000 -0.188982
B -0.188982  1.000000

```

3.2 Summary Functions

3.2.1 numeric_summary

`PyAudit.basics.numeric_summary(df_in, output_dir, top_n=4, deciles=False)`
generate statistical summary for numerical DataFrame

Parameters

- **df_in** – input pandas DataFrame
- **deciles** – flag for percentiles style

Returns statistical summary for numerical data

Author Wenqiang Feng and Ming Chen

Email von198@gmail.com

```

>>> d = {'A': [1, 0, None, 3],
>>>       'B': [1, 0, 0, 0],
>>>       'C': ['a', None, 'c', 'd']}
>>> # create DataFrame
>>> df = pd.DataFrame(d)
>>> print(numeric_summary(df))
  feature data_type  min_digits  ...  zero_rate  pos_rate  neg_
↪ rate
  A      A    float64          3  ...    0.333333  0.666667  ↪
↪ 0.0
  B      B    int64          3  ...    0.750000  0.250000  ↪
↪ 0.0

```

3.2.2 category_summary

`PyAudit.basics.category_summary(df_in, output_dir, top_n=4, deciles=False)`
generate statistical summary for numerical DataFrame

Parameters

- **df_in** – input pandas DataFrame
- **deciles** – flag for percentiles style

Returns statistical summary for numerical data

Author Wenqiang Feng and Ming Chen

Email von198@gmail.com

```
>>> d = {'A': [1, 0, None, 3],
>>>       'B': [1, 0, 0, 0],
>>>       'C': ['a', None, 'c', 'd']}
>>> # create DataFrame
>>> df = pd.DataFrame(d)
>>> print(numeric_summary(df))
   feature data_type  min_digits  ...  top_values  top_freqs  _
->missing_rate
   C      C      object          1  ...   [a, d, c]  [1, 1, 1]  _
->      0.25
```

3.3 Auditing Function

3.3.1 auditing

AUDITING DEMOS

The following demos are designed to show how to use `PyAudit` to audit `pd.DataFrame`.

4.1 Auditing in one function

For example:

Result:

	Age	Sex	ChestPain	RestBP	Chol	...	Oldpeak	Slope	Ca	
→	Thal	AHD								
0	63	True	typical	145	233	...	2.3	3	0.0	
→	fixed	No								
1	67	True	asymptomatic	160	286	...	1.5	2	3.0	
→	normal	Yes								
2	67	True	asymptomatic	120	229	...	2.6	2	2.0	
→	reversable	Yes								
3	37	True	nonanginal	130	250	...	3.5	3	0.0	
→	normal	No								
4	41	False	nontypical	130	204	...	1.4	1	0.0	
→	normal	No								
[5 rows x 14 columns]										
	feature	data_type	min_digits	...	zero_rate	pos_rate	neg_			
→	rate									
Age	Age	int64	4	...	0.000000	1.000000				
→	0.0									
RestBP	RestBP	int64	4	...	0.000000	1.000000				
→	0.0									
Chol	Chol	int64	5	...	0.000000	1.000000				
→	0.0									
Fbs	Fbs	int64	3	...	0.851485	0.148515				
→	0.0									

(continues on next page)

(continued from previous page)

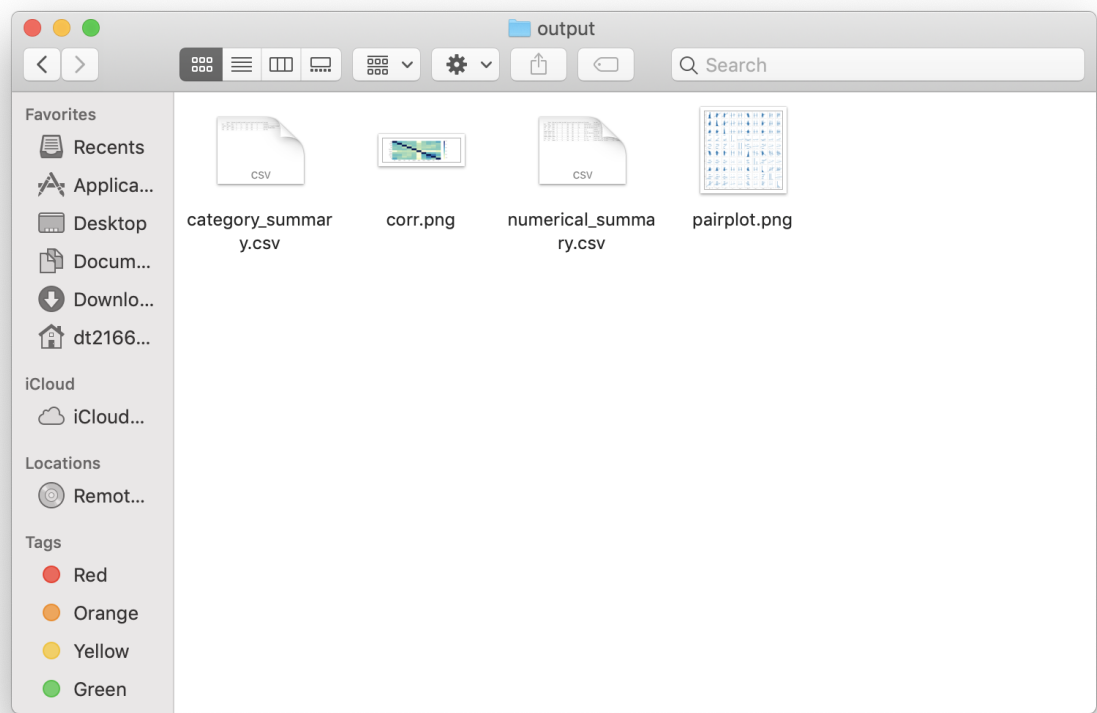
RestECG	RestECG	int64	3	...	0.498350	0.501650	└
→0.0							
MaxHR	MaxHR	int64	4	...	0.000000	1.000000	└
→0.0							
ExAng	ExAng	int64	3	...	0.673267	0.326733	└
→0.0							
Oldpeak	Oldpeak	float64	3	...	0.326733	0.673267	└
→0.0							
Slope	Slope	int64	3	...	0.000000	1.000000	└
→0.0							
Ca	Ca	float64	3	...	0.588629	0.411371	└
→0.0							
[10 rows x 21 columns]							
	feature	data_type	...		top_freqs	missing_rate	
Sex	Sex	bool	...		[206, 97]	0.000000	
ChestPain	ChestPain	object	...	[144, 86, 50, 23]		0.000000	
Thal	Thal	object	...	[166, 117, 18]		0.006601	
AHD	AHD	object	...	[164, 139]		0.000000	
[4 rows x 10 columns]							
	Age	RestBP	Chol	...	Oldpeak	Slope	└
→Ca							
Age	1.000000	0.284946	0.208950	...	0.203805	0.161770	0.
→362605							
RestBP	0.284946	1.000000	0.130120	...	0.189171	0.117382	0.
→098773							
Chol	0.208950	0.130120	1.000000	...	0.046564	-0.004062	0.
→119000							
Fbs	0.118530	0.175340	0.009841	...	0.005747	0.059894	0.
→145478							
RestECG	0.148868	0.146560	0.171043	...	0.114133	0.133946	0.
→128343							
MaxHR	-0.393806	-0.045351	-0.003432	...	-0.343085	-0.385601	-0.
→264246							
ExAng	0.091661	0.064762	0.061310	...	0.288223	0.257748	0.
→145570							
Oldpeak	0.203805	0.189171	0.046564	...	1.000000	0.577537	0.
→295832							
Slope	0.161770	0.117382	-0.004062	...	0.577537	1.000000	0.
→110119							
Ca	0.362605	0.098773	0.119000	...	0.295832	0.110119	1.
→000000							
[10 rows x 10 columns]							

(continues on next page)

(continued from previous page)

```
Process finished with exit code 0
```

and



4.2 Auditing function by function

For example:

Result:

	Age	Sex	ChestPain	RestBP	Chol	...	Oldpeak	Slope	Ca	
→	Thal	AHD								
0	63	True	typical	145	233	...	2.3	3	0.0	
→	fixed	No								
1	67	True	asymptomatic	160	286	...	1.5	2	3.0	
→	normal	Yes								

(continues on next page)

(continued from previous page)

2	67	True	asymptomatic	120	229	...	2.6	2	2.0	└
↪	reversible	Yes								
3	37	True	nonanginal	130	250	...	3.5	3	0.0	└
↪	normal	No								
4	41	False	nontypical	130	204	...	1.4	1	0.0	└
↪	normal	No								
[5 rows x 14 columns]										
	feature	data_type	min_digits	...	zero_rate	pos_rate	neg_			
↪	rate									
Age	Age	int64	4	...	0.000000	1.000000				└
↪	0.0									
RestBP	RestBP	int64	4	...	0.000000	1.000000				└
↪	0.0									
Chol	Chol	int64	5	...	0.000000	1.000000				└
↪	0.0									
Fbs	Fbs	int64	3	...	0.851485	0.148515				└
↪	0.0									
RestECG	RestECG	int64	3	...	0.498350	0.501650				└
↪	0.0									
MaxHR	MaxHR	int64	4	...	0.000000	1.000000				└
↪	0.0									
ExAng	ExAng	int64	3	...	0.673267	0.326733				└
↪	0.0									
Oldpeak	Oldpeak	float64	3	...	0.326733	0.673267				└
↪	0.0									
Slope	Slope	int64	3	...	0.000000	1.000000				└
↪	0.0									
Ca	Ca	float64	3	...	0.588629	0.411371				└
↪	0.0									
[10 rows x 21 columns]										
	feature	data_type	...	top_freqs	missing_rate					
Sex	Sex	bool	...	[206, 97]	0.000000					
ChestPain	ChestPain	object	...	[144, 86, 50, 23]	0.000000					
Thal	Thal	object	...	[166, 117, 18]	0.006601					
AHD	AHD	object	...	[164, 139]	0.000000					
[4 rows x 10 columns]										
	Age	RestBP	Chol	...	Oldpeak	Slope				└
↪	Ca									
Age	1.000000	0.284946	0.208950	...	0.203805	0.161770	0.			
↪	362605									
RestBP	0.284946	1.000000	0.130120	...	0.189171	0.117382	0.			
↪	098773									

(continues on next page)

(continued from previous page)

```

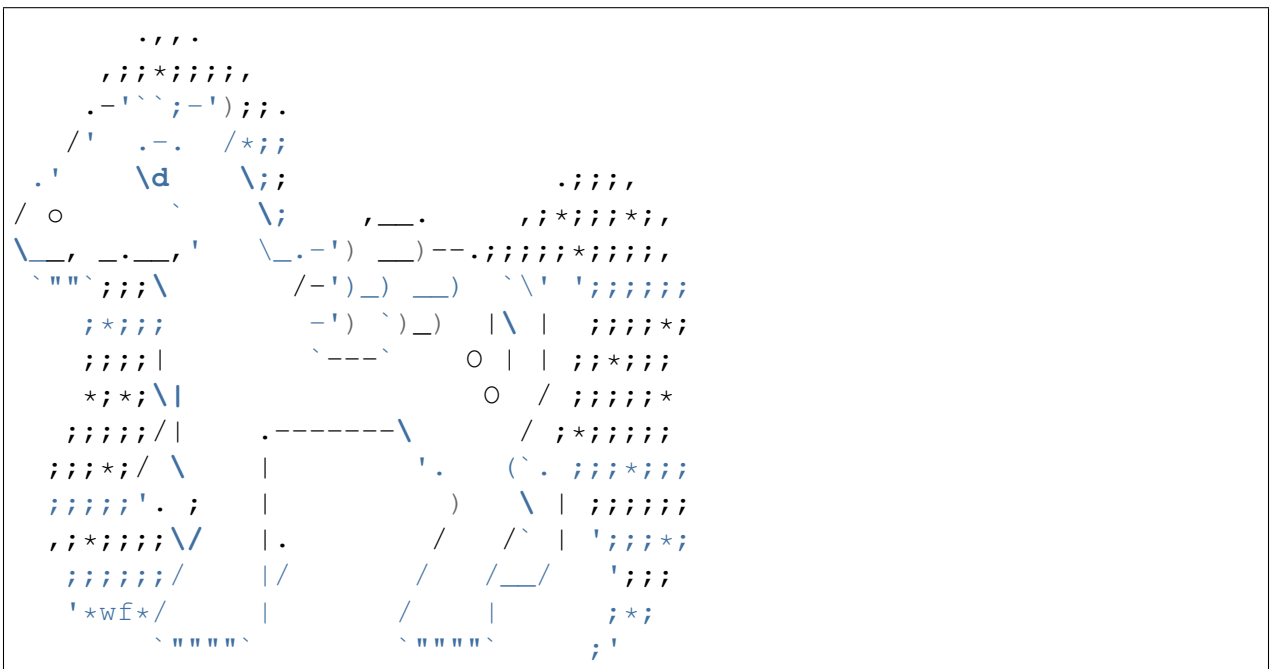
Chol      0.208950  0.130120  1.000000  ...  0.046564 -0.004062  0.
↪119000
Fbs       0.118530  0.175340  0.009841  ...  0.005747  0.059894  0.
↪145478
RestECG   0.148868  0.146560  0.171043  ...  0.114133  0.133946  0.
↪128343
MaxHR     -0.393806 -0.045351 -0.003432  ... -0.343085 -0.385601 -0.
↪264246
ExAng     0.091661  0.064762  0.061310  ...  0.288223  0.257748  0.
↪145570
Oldpeak   0.203805  0.189171  0.046564  ...  1.000000  0.577537  0.
↪295832
Slope     0.161770  0.117382 -0.004062  ...  0.577537  1.000000  0.
↪110119
Ca        0.362605  0.098773  0.119000  ...  0.295832  0.110119  1.
↪000000

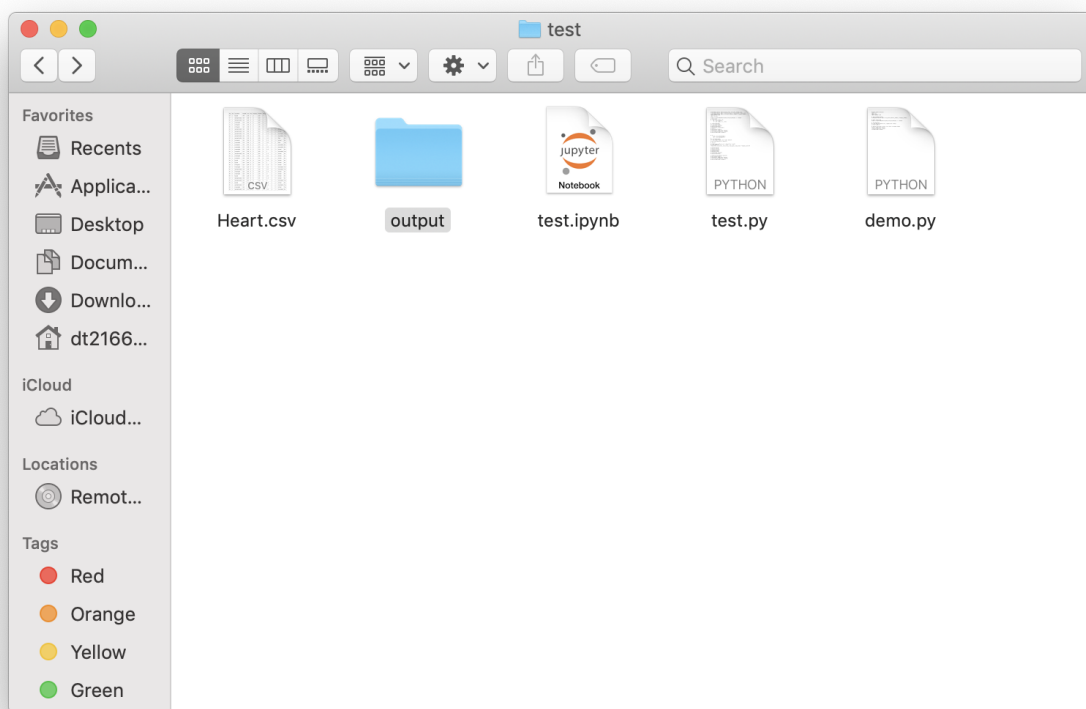
```

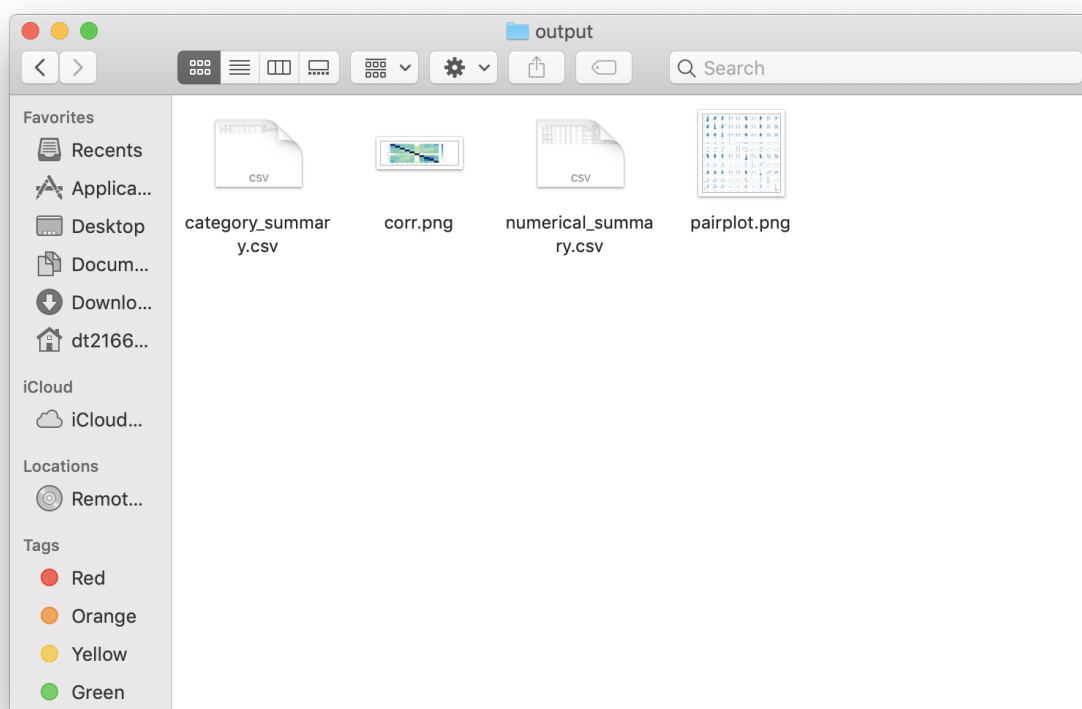
```
[10 rows x 10 columns]
```

Process finished with `exit` code 0

and







MAIN REFERENCE

BIBLIOGRAPHY

[PyAudit] Wenqiang Feng and Ming Chen. [Python Data Audit Library API](#), 2019.

PYTHON MODULE INDEX

p

`PyAudit.basics`, 9

INDEX

D

`dtypes_class()` (*in module PyAudit basics*), 9

P

`PyAudit.basics` (*module*), 9